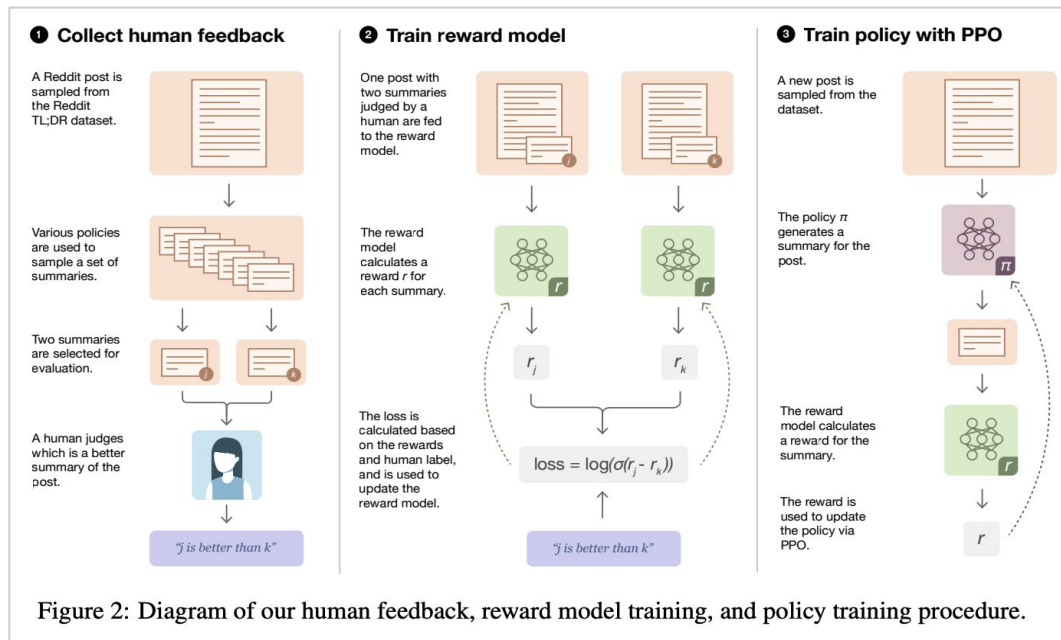


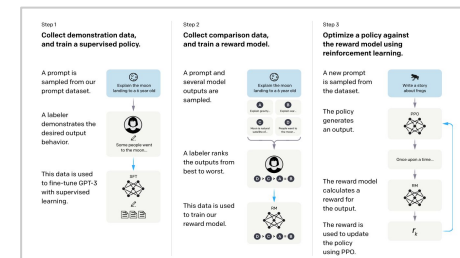
PPO

Let's get confused together :)





2020



2022

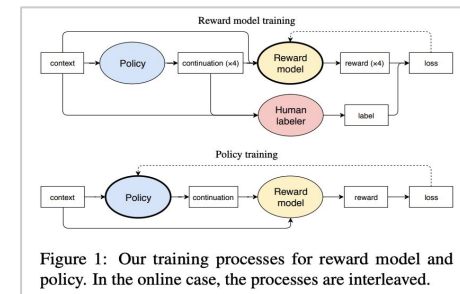
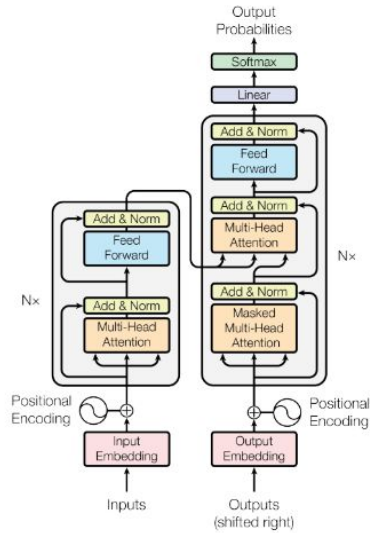


Figure 1: Our training processes for reward model and policy. In the online case, the processes are interleaved.

2019

Able to store massive
amounts of data



Efficient knowledge
extraction





May 19, 2025

Does Reinforcement Learning Really Incentivize Reasoning Capacity in LLMs Beyond the Base Model?

Yang Yue^{1,†}, Zhiqi Chen^{1,*}, Rui Lu¹, Andrew Zhao¹, Zhaokai Wang², Yang Yue¹, Shiji Song¹, and Gao Huang^{1,✉}

¹LeapLab, Tsinghua University ²Shanghai Jiao Tong University

* Equal Contribution [†] Project Lead [✉] Corresponding Author

Reinforcement Learning with Verifiable Rewards (RLVR) has recently demonstrated notable success in enhancing the reasoning performance of large language models (LLMs), particularly in mathematics and programming tasks. It is widely believed that, similar to how traditional RL helps agents to explore and learn new strategies, RLVR enables LLMs to continuously self-improve, thus acquiring novel reasoning abilities that exceed the capacity of the corresponding base models. In this study, we take a critical look at *the current state of RLVR* by systematically probing the reasoning capability boundaries of RLVR-trained LLMs across various model families, RL algorithms, and math/coding/visual reasoning benchmarks, using pass@k at large k values as the evaluation metric. While RLVR improves sampling efficiency towards correct paths, we surprisingly find that current training does *not* elicit fundamentally new reasoning patterns. We observe that while RLVR-trained models outperform their base models at smaller values of k (e.g., $k=1$), base models achieve higher pass@k score when k is large. Moreover, we observe that the reasoning capability boundary of LLMs often narrows as RLVR training progresses. Further coverage and perplexity analysis shows that the reasoning paths generated by RLVR models are already included in the base models' sampling distribution, suggesting that their reasoning abilities originate from and are *bounded* by the base model. From this perspective, treating the base model as an upper bound, our quantitative analysis shows that six popular RLVR algorithms perform similarly and remain far from optimal in fully leveraging the potential of the base model. In contrast, we find that distillation can introduce new reasoning patterns from the teacher and genuinely expand the model's reasoning capabilities. Taken together, our findings suggest that current RLVR methods have not fully realized the potential of RL to elicit genuinely novel reasoning abilities in LLMs. This underscores the need for improved RL paradigms—such as continual scaling and multi-turn agent-environment interaction—to unlock this potential.


Project Page: <https://limit-of-RLVR.github.io>

1. Introduction




The development of reasoning-centric large language models (LLMs), such as OpenAI-o1 (Jaech et al., 2024), DeepSeek-R1 (Guo et al., 2025), and Kimi-1.5 (Team et al., 2025), has significantly advanced the frontier of LLM capabilities, particularly in solving complex logical tasks involving mathematics and programming. In contrast to traditional instruction-tuned approaches that rely on human-curated annotations (Achiam et al., 2023; Grattafiori et al., 2024), the key driver behind this leap forward is large-scale Reinforcement Learning with Verifiable Rewards (RLVR) (Lambert et al., 2024; Guo et al.,

Correspond to: {lc-yz2, zq-chen25}@mails.tsinghua.edu.cn, gao.huang@tsinghua.edu.cn.

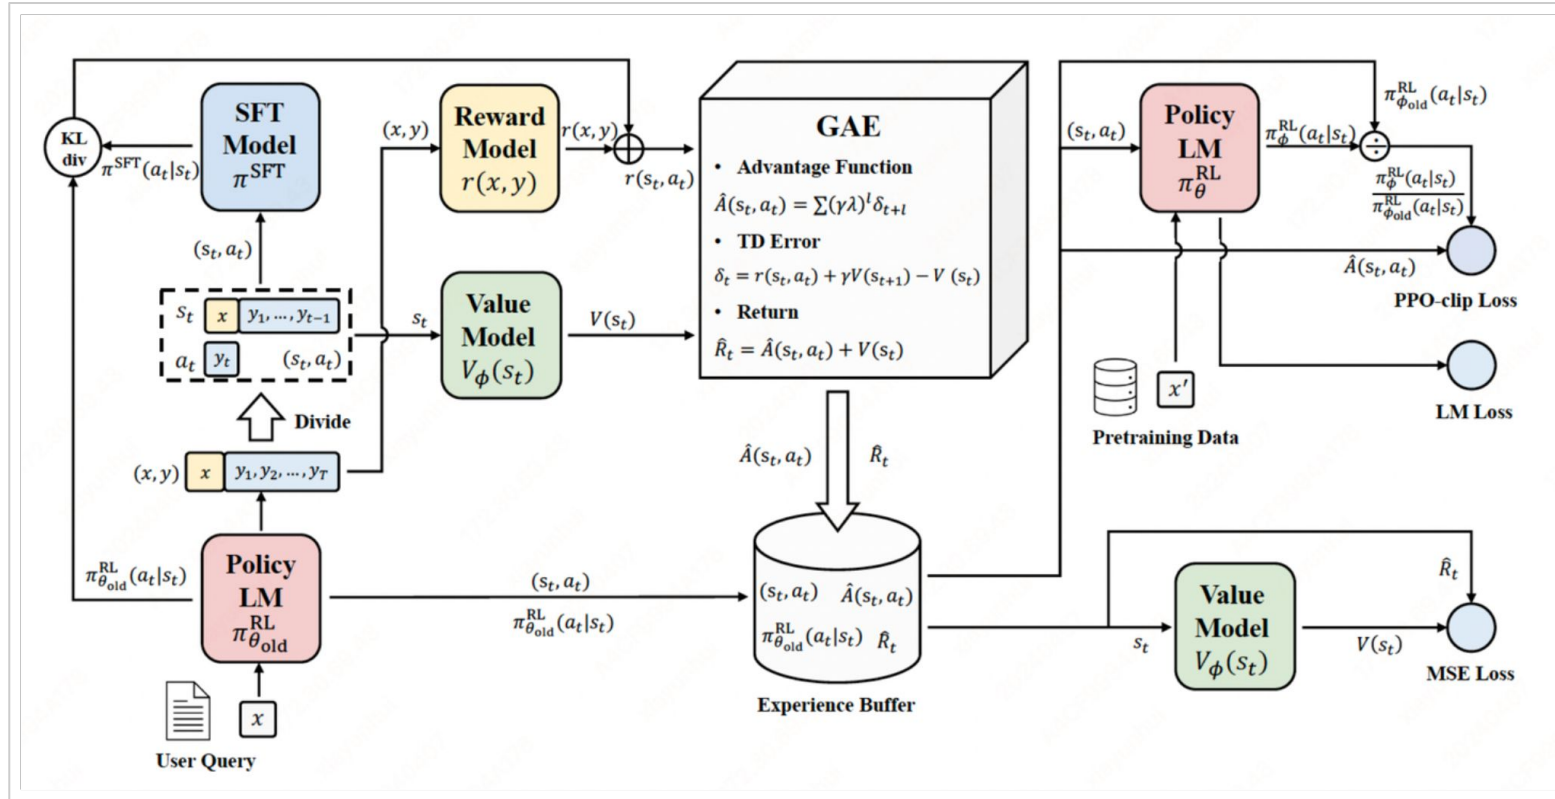
Minerva



Olympiad





Overview



Bradley-Terry

$$\mathcal{L}(\theta) = -\mathbb{E}_{(x, y_w, y_l)} \left[\log \sigma(r_\theta(x, y_w) - r_\theta(x, y_l)) \right] = -\mathbb{E}_{(x, y_w, y_l)} \log \left[\frac{e^{r_\theta(x, y_w)}}{e^{r_\theta(x, y_w)} + e^{r_\theta(x, y_l)}} \right]$$



$$s_t = (x, y_{1:t-1}) \quad (1)$$

$$a_t = y_t \quad (2)$$

$$r_t = -\beta D_{\text{KL}}(\pi_\theta(\cdot \mid s_t) \parallel \pi_{\text{ref}}(\cdot \mid s_t)) + \mathbb{1}_{\{t=T\}} r_{\text{score}}(x, y). \quad (3)$$

$$D_{\text{KL}}^{(t)} \approx \log \pi_\theta(a_t \mid s_t) - \log \pi_{\text{ref}}(a_t \mid s_t). \quad (4)$$

$$\delta_t = r_t + \gamma V_\phi(s_{t+1}) - V_\phi(s_t). \quad (5)$$

$$\hat{A}_t = \sum_{l=0}^{T-t-1} (\gamma\lambda)^l \delta_{t+l}. \quad (6)$$

$$\hat{R}_t = \hat{A}_t + V_\phi(s_t). \quad (7)$$

$$r_t(\theta) = \frac{\pi_\theta(a_t \mid s_t)}{\pi_{\text{old}}(a_t \mid s_t)}. \quad (8)$$

$$r_t(\theta) = \frac{\pi_\theta(a_t \mid s_t)}{\pi_{\text{old}}(a_t \mid s_t)}. \quad (8)$$

$$J_{\text{clip}}(\theta) = \mathbb{E}_t \left[\min \left(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right) \right]. \quad (9)$$

$$L_{\text{policy}}(\theta) = -J_{\text{clip}}(\theta). \quad (10)$$


$$L_{\text{value}}(\phi) = \mathbb{E}_t \left[(V_\phi(s_t) - \hat{R}_t)^2 \right]. \quad (11)$$

$$L_{\text{ent}}(\theta) = -\mathbb{E}_t [\mathcal{H}(\pi_\theta(\cdot \mid s_t))]. \quad (12)$$

$$L_{\text{ptx}}(\theta) = -\mathbb{E}_{(x,y) \sim \mathcal{D}_{\text{ptx}}} \sum_{t=1}^T \log \pi_\theta(y_t \mid x, y_{1:t-1}). \quad (13)$$

$$L_{\text{total}}(\theta, \phi) = L_{\text{policy}}(\theta) + c_v L_{\text{value}}(\phi) + c_{\text{ent}} L_{\text{ent}}(\theta) + c_{\text{ptx}} L_{\text{ptx}}(\theta). \quad (14)$$

Policy & Friends



yes, everything is calculated at the token level, it is a form of reward shaping

this is calculated for every single token, correct?

only the last token gets the extra reward score, previous tokens get only the approximate KL divergence



what is this?

$$r_t = -\beta D_{\text{KL}}(\pi_{\theta}(\cdot | s_t) \| \pi_{\text{ref}}(\cdot | s_t)) + \mathbb{1}_{\{t=T\}} r_{\text{score}}(x, y). \quad (3)$$


$$D_{\text{KL}}^{(t)} \approx \log \pi_{\theta}(a_t | s_t) - \log \pi_{\text{ref}}(a_t | s_t). \quad (4)$$

why an approximate?


approximated via sampling, because the true KL divergence would require the all distribution or all the possible actions



$$D_{\text{KL}}^{(t)} \approx \log \pi_{\theta}(a_t | s_t) - \log \pi_{\text{ref}}(a_t | s_t)$$




Ref




Policy

$$r_t(\theta) = \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\text{old}}(a_t | s_t)}$$



Policy



Old

$$r_t(\theta) = \frac{\pi_\theta(a_t \mid s_t)}{\pi_{\text{old}}(a_t \mid s_t)}. \quad (8)$$

$$J_{\text{clip}}(\theta) = \mathbb{E}_t \left[\min \left(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right) \right]. \quad (9)$$

$$L_{\text{policy}}(\theta) = -J_{\text{clip}}(\theta). \quad (10)$$

$$L_{\text{value}}(\phi) = \mathbb{E}_t \left[(V_\phi(s_t) - \hat{R}_t)^2 \right]. \quad (11)$$

$$L_{\text{ent}}(\theta) = -\mathbb{E}_t [\mathcal{H}(\pi_\theta(\cdot \mid s_t))]. \quad (12)$$

$$L_{\text{ptx}}(\theta) = -\mathbb{E}_{(x,y) \sim \mathcal{D}_{\text{ptx}}} \sum_{t=1}^T \log \pi_\theta(y_t \mid x, y_{1:t-1}). \quad (13)$$

$$L_{\text{total}}(\theta, \phi) = L_{\text{policy}}(\theta) + c_v L_{\text{value}}(\phi) + c_{\text{ent}} L_{\text{ent}}(\theta) + c_{\text{ptx}} L_{\text{ptx}}(\theta). \quad (14)$$

$$\text{clip}(x, a, b) = \begin{cases} a, & x < a, \\ x, & a \leq x \leq b, \\ b, & x > b. \end{cases}$$